


# Round-Robin (RR) Scheduling

- Each process gets a small unit of CPU time (*time quantum*), usually 10~100 ms
- After TQ elapsed, process is preempted and added to the end of the ready queue
- Performance
  - TQ large → FIFO
  - TQ small → (context switch) overhead increases

# RR Scheduling (TQ = 20)

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

■ The Gantt chart is:



■ Typically, higher average turnaround than SJF, but better *response*.