


Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
- A process with shortest burst length gets the CPU first
- SJF provides the minimum average waiting time (optimal!)
- Two schemes
 - **Non-preemptive** – once CPU given to a process, it cannot be preempted until its completion
 - **Preemptive** – if a new process arrives with shorter burst length, preemption happens

Non-Preemptive SJF Example

| <u>Process</u> | <u>Arrival Time</u> | <u>Burst Time</u> |
|----------------|---------------------|-------------------|
| P1 | 0 | 7 |
| P2 | 2 | 4 |
| P3 | 4 | 1 |
| P4 | 5 | 4 |

Ready queue: t=12



Wait time = completion time – arrival time – run time (burst time)


$$\text{AWT} = [(7-0-7)+(12-2-4)+(8-4-1)+(16-5-4)]/4 = (0+6+3+7)/4 = 4$$

Response Time: P1=0, P2=6, P3=3, P4=7

Preemptive SJF Example

| <u>Process</u> | <u>Arrival Time</u> | <u>Burst Time</u> |
|----------------|---------------------|-------------------|
| P1 | 0 | 7 |
| P2 | 2 | 4 |
| P3 | 4 | 1 |
| P4 | 5 | 4 |

Ready queue: t=11



Wait time = completion time – arrival time – run time (burst time)

$$\text{AWT} = [(16-0-7)+(7-2-4)+(5-4-1)+(11-5-4)]/4 = (9+1+0+2)/4 = \mathbf{3}$$

Response Time: P1=0, P2=0, P3=0, P4=2

Approximate Shortest-Job-First (SJF)

- SJF difficulty: no way to know length of the next CPU burst
- **Approximate SJF**: the next burst can be predicted as an **exponential average** of the measured length of previous CPU bursts

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \pi_n$$

new one \rightarrow t_n \leftarrow history π_n

Commonly,

$$\begin{aligned} &= \alpha t_n + (1 - \alpha) \alpha t_{n-1} + (1 - \alpha)^2 \alpha t_{n-2} + \dots \\ \alpha = 1/2 &\longrightarrow \\ &= \left(\frac{1}{2}\right) t_n + \left(\frac{1}{2}\right)^2 t_{n-1} + \left(\frac{1}{2}\right)^3 t_{n-2} + \dots \end{aligned}$$